


# VPCs & Multiple EC2 Instances

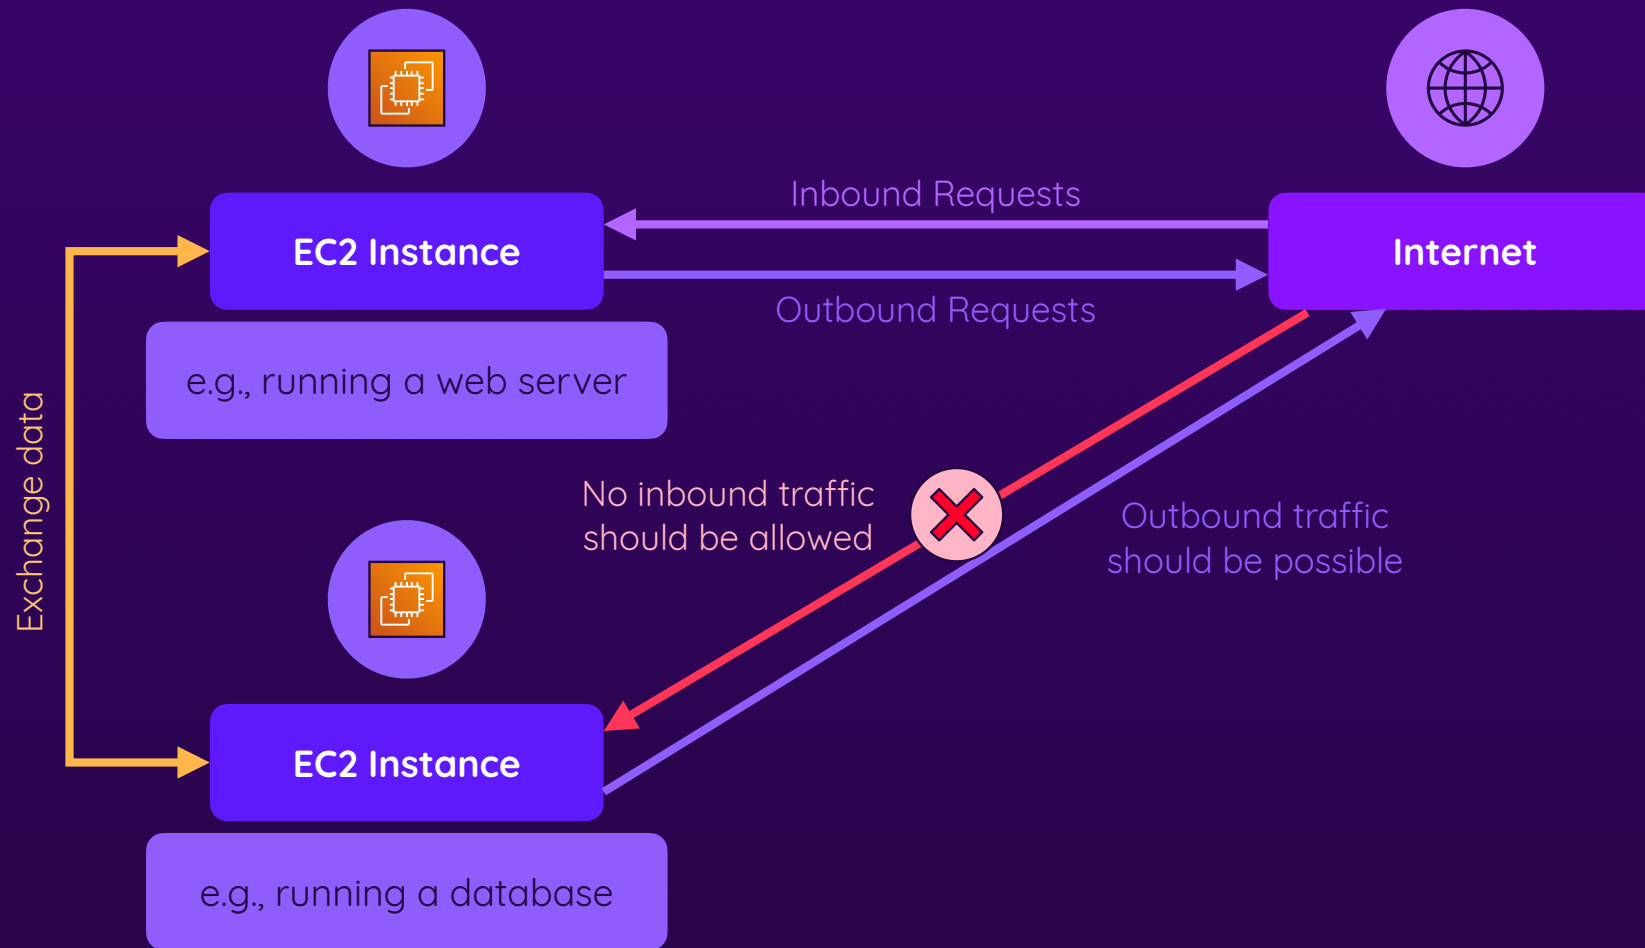
Managing your own network in the cloud

- ▶ Understanding VPCs
- ▶ Private vs Public Instances
- ▶ Managing Network Requests

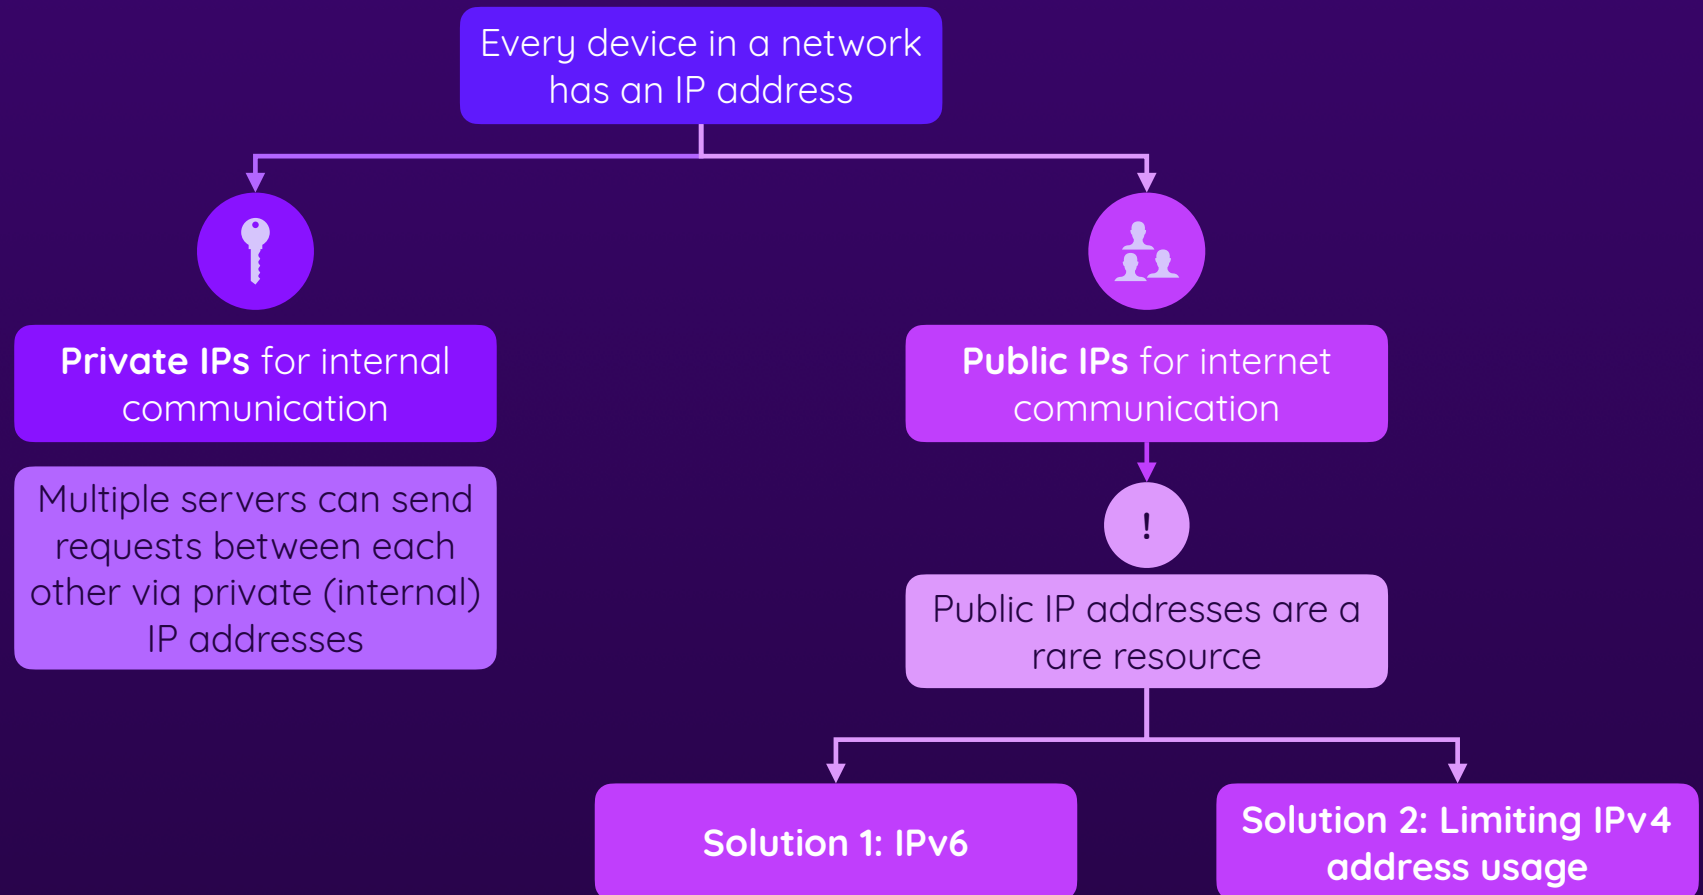
# A Simple Setup



# A More Realistic Setup




# Public vs Private IP Addresses



# Understanding IP (IPv4) Addresses

An IP address is a 32-bit number

172 . 31 . 0 . 0



4 x 8-bit

This is just a notation thing  
though (for human readability)

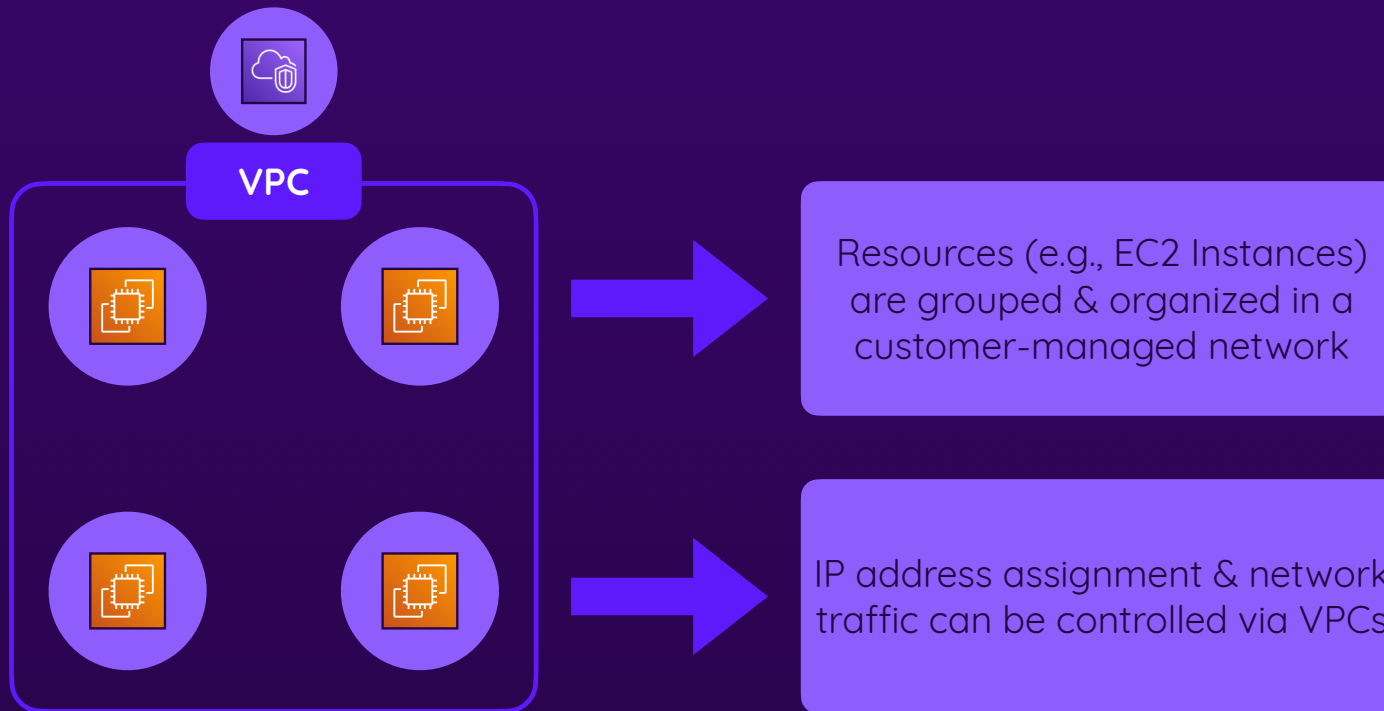
# IPv4 Addresses Are A Rare Resource

Less than 4.3bn available IPv4 addresses




Not enough for all the devices  
(with internet access) we have  
around the world

# Introducing Virtual Private Clouds (VPCs)



# VPCs & Subnets




You actually control network request settings on subnet-level

This allows you to make subnets “**private**” (only internal requests) or “**public**” (internet requests are possible)



# Public vs Private Subnets



# What About Security Groups?



## Security Group

Firewalls, directly attached to EC2 instances

Technically, requests still reach the instances

Security groups just allow which requests to allow or block



## Private Subnets

Not directly attached to instances

Instead: Contain multiple EC2 instances

Technical isolation from the internet: Requests technically don't reach the instances

# Public vs Private Subnets



## Public Subnet

Associated with a route table that has an internet gateway route

Instances can communicate with each other **AND the internet**



## Private Subnet

Associated with a route table that has **no internet gateway route**

Instances can communicate with each other only



To still allow for outgoing internet access, a **NAT gateway** can be used

# Understanding CIDR IP Ranges

An IP address is a 32-bit number

172 . 31 . 0 . 0 / 16



Defines how many  
bits are fixed

4 x 8-bit

This is just a notation thing though (for  
human readability)

# Understanding CIDR IP Ranges

172.31.0.0 / 16

16 bits are fixed

Range

172.31.0.0



172.31.255.255

65,536 available  
addresses

172.31.0.0 / 24

24 bits are fixed

Range

172.31.0.0



172.31.0.255

256 available  
addresses

0.0.0.0 / 0

0 bits are fixed

Range

Unlimited

All possible IP  
addresses

A higher /X number implies less available IP addresses in the range

# Elastic IPs

**Automatically assigned IPs** (by subnet) will change  
when instances are stopped / restarted

You can't control which public IP address gets  
assigned to an instance



**Elastic IPs are managed & assigned by you**

Elastic IPs don't change and can be re-assigned

# Always Use Elastic IPs?

**Automatically assigned IPs** (by subnet) will change when instances are stopped / restarted

You can't control which public IP address gets assigned to an instance



**Elastic IPs are managed & assigned by you**

Elastic IPs don't change and can be re-assigned

**Elastic IPs should be used with care**

Scarce resource: You can only have a few EIPs per region & account

Unused EIPs incur charges

**There often are better alternatives**

e.g., use DNS for exposing applications / websites to the world

e.g., use application integration services (like SQS) for connecting workloads

# Security Groups & Network ACLs

## Security Group

Preferred

Firewall for a single EC2 instance

Checks incoming / outgoing requests and conditionally blocks or allows them

Stateful: Responses are always allowed (if the request passed)

Multiple instances can have different security groups

Security groups can be re-used for multiple instances

## Network ACL (Access Control List)

Firewall for entire subnets

Checks incoming / outgoing requests and conditionally blocks or allows them

Stateless: Requests & responses are decoupled

One NACL can be associated with multiple subnets

## Private / Public Subnets

Defines technical connectivity

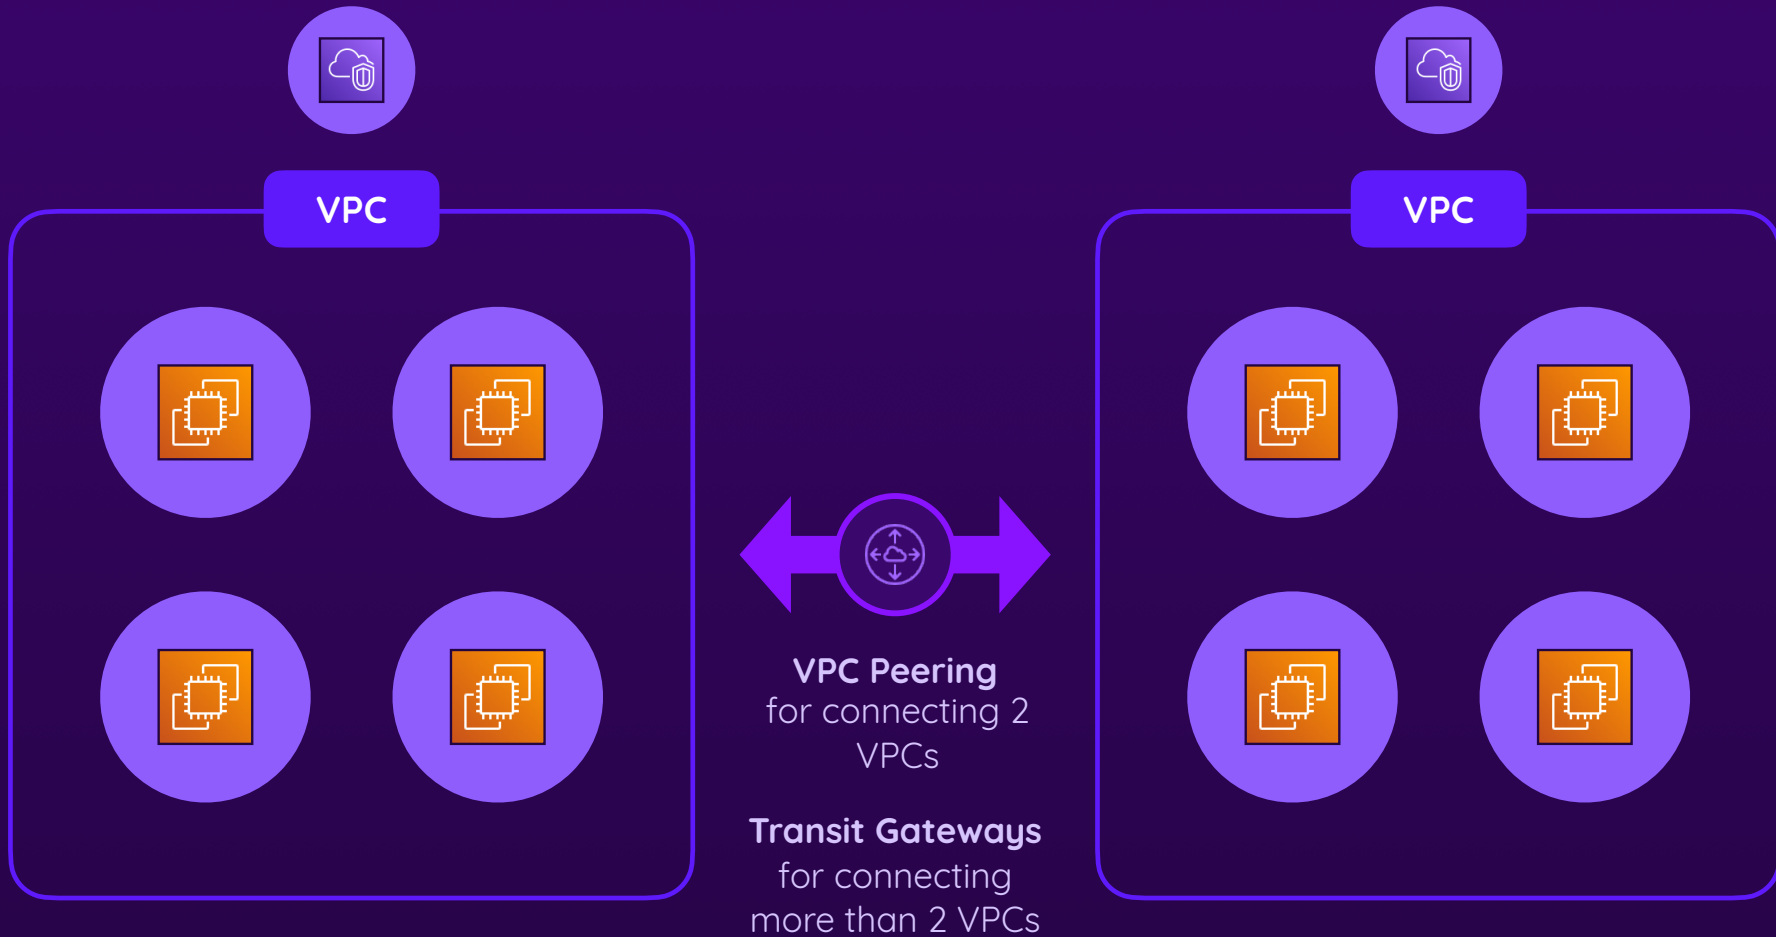
No internet access without internet gateway (incoming + outgoing) or NAT gateway (outgoing)

Does not control any requests or responses

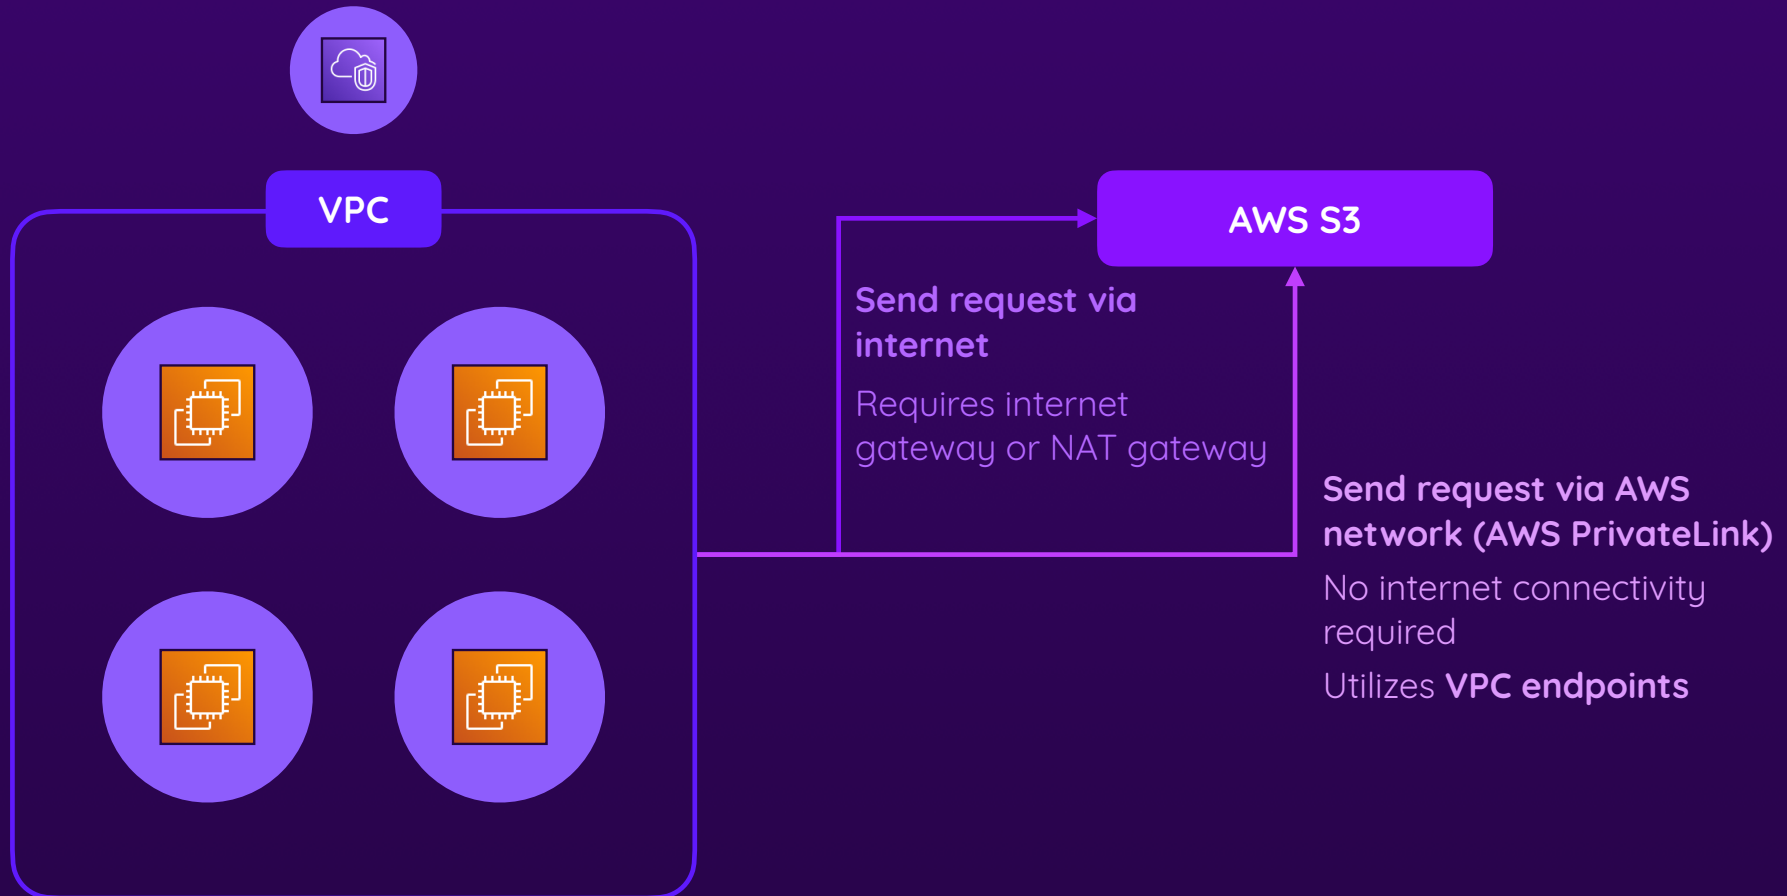
Multiple instances can be in the same subnet



# VPC Peering & Transit Gateways



# VPC Endpoints & AWS PrivateLink



# Summary



## VPCs (Virtual Private Cloud)

Your own network in the cloud  
(for grouping EC2 instances)

A VPC contains at least two  
subnets & one route table

Subnets can be “**public**” or  
“**private**”

Route table settings control  
subnet “visibility”



## Network Management

Every VPC has an IP CIDR block  
assigned (range of IPs)

Subnets get parts of the VPC  
CIDR block assigned

EC2 instances receive auto-  
assigned public and private IPs

Elastic IPs can be used for fixed IP  
addresses

VPC peering or transit gateways  
can connect VPCs



## Request Management

Internet gateways allow for two-  
way internet access

NAT gateways enable outgoing  
internet requests

NACLs allow or deny requests on  
subnet-level

Endpoints (PrivateLink) connect  
AWS services to VPCs